

Australian Rain Predictor

Andrew White

1/10/2020

Introduction

This project looks at a dataset pulled from Kaggle that contains the weather data from 2007-2017 for Australia from 44 different cities. The goal of the project is to predict if it will rain tomorrow using the other weather variables. There is a lot of practical use for weather modeling, especially in Australia as they battle massive and devastating wildfires, since predicting the weather allows the firefighters to strategize where to defend next. Using data visualization, I will determine the best predictors then use them to build the model. I will use a decision tree model to make a prediction then a random forest in an attempt to improve the prediction. I will use Accuracy, Sensitivity, Specificity, and Balanced Accuracy to measure the success of each model. I am using the packages- tidyverse, caret, lubridate, rpart, and randomForest.

Data loading and cleaning

The first step is to load all the required packages

```
if(!require(tidyverse)) install.packages("tidyverse", repos = "http://cran.us.r-project.org")
if(!require(caret)) install.packages("caret", repos = "http://cran.us.r-project.org")
if(!require(lubridate)) install.packages("lubridate", repos = "http://cran.us.r-project.org")
if(!require(rpart)) install.packages("rpart", repos = "http://cran.us.r-project.org")
if(!require(rpart.plot)) install.packages("rpart.plot", repos = "http://cran.us.r-project.org")
if(!require(randomForest)) install.packages("randomForest", repos = "http://cran.us.r-project.org")
```

Next load the dataset

```
url<-"https://github.com/adwhite3/Edx-Final/blob/master/weatherAUS.csv?raw=TRUE"
data<-read_csv(url)
```

```
## Parsed with column specification:
## cols(
##   .default = col_double(),
##   Date = col_date(format = ""),
##   Location = col_character(),
##   Evaporation = col_logical(),
##   Sunshine = col_logical(),
##   WindGustDir = col_character(),
##   WindDir9am = col_character(),
##   WindDir3pm = col_character(),
##   RainToday = col_character(),
##   RainTomorrow = col_character()
## )

## See spec(...) for full column specifications.
```

I will be making a classification tree to predict if there will or will not be precipitation tomorrow in Australia. Data ranges from 2007-2017. First I remove RISK_MM column because it is future information that will interfere with the prediction.

```
data$RISK_MM<-NULL
```

Now I can begin looking at data

```
summary(data)
```

```
##      Date      Location      MinTemp      MaxTemp
## Min.   :2007-11-01  Length:142193  Min.   :-8.50  Min.   :-4.80
## 1st Qu.:2011-01-06  Class :character 1st Qu.: 7.60  1st Qu.:17.90
## Median :2013-05-27  Mode  :character Median :12.00  Median :22.60
## Mean   :2013-04-01      Mean   :12.19  Mean   :23.23
## 3rd Qu.:2015-06-12      3rd Qu.:16.80  3rd Qu.:28.20
## Max.   :2017-06-25      Max.   :33.90  Max.   :48.10
##                      NA's   :637    NA's   :322
##      Rainfall      Evaporation      Sunshine      WindGustDir
## Min.   : 0.00      Mode :logical  Mode :logical  Length:142193
## 1st Qu.: 0.00      FALSE:240    FALSE:2308    Class :character
## Median : 0.00      TRUE :1563   TRUE :326     Mode  :character
## Mean   : 2.35      NA's :140390 NA's :139559
## 3rd Qu.: 0.80
## Max.   :371.00
## NA's   :1406
## WindGustSpeed      WindDir9am      WindDir3pm      WindSpeed9am
## Min.   : 6.00      Length:142193  Length:142193  Min.   : 0
## 1st Qu.:31.00      Class :character  Class :character 1st Qu.: 7
## Median :39.00      Mode  :character  Mode  :character Median :13
## Mean   :39.98
## 3rd Qu.:48.00
## Max.   :135.00
## NA's   :9270
##                      NA's   :1348
## WindSpeed3pm      Humidity9am      Humidity3pm      Pressure9am
## Min.   : 0.00      Min.   : 0.00  Min.   : 0.00  Min.   : 980.5
## 1st Qu.:13.00      1st Qu.: 57.00 1st Qu.: 37.00 1st Qu.:1012.9
## Median :19.00      Median : 70.00 Median : 52.00 Median :1017.6
## Mean   :18.64      Mean   : 68.84 Mean   : 51.48 Mean   :1017.7
## 3rd Qu.:24.00      3rd Qu.: 83.00 3rd Qu.: 66.00 3rd Qu.:1022.4
## Max.   :87.00      Max.   :100.00 Max.   :100.00 Max.   :1041.0
## NA's   :2630      NA's   :1774   NA's   :3610   NA's   :14014
## Pressure3pm      Cloud9am      Cloud3pm      Temp9am
## Min.   : 977.1      Min.   :0.00    Min.   :0.0     Min.   : -7.20
## 1st Qu.:1010.4      1st Qu.:1.00    1st Qu.:2.0     1st Qu.:12.30
## Median :1015.2      Median :5.00    Median :5.0     Median :16.70
## Mean   :1015.3      Mean   :4.44    Mean   :4.5     Mean   :16.99
## 3rd Qu.:1020.0      3rd Qu.:7.00    3rd Qu.:7.0     3rd Qu.:21.60
## Max.   :1039.6      Max.   :9.00    Max.   :9.0     Max.   :40.20
## NA's   :13981      NA's   :53657   NA's   :57094   NA's   :904
## Temp3pm      RainToday      RainTomorrow
## Min.   : -5.40  Length:142193  Length:142193
## 1st Qu.:16.60  Class :character  Class :character
```

```
## Median :21.10   Mode  :character   Mode  :character
## Mean    :21.69
## 3rd Qu. :26.40
## Max.    :46.70
## NA's    :2726
```

```
head(data)
```

```
## # A tibble: 6 x 23
##   Date      Location MinTemp MaxTemp Rainfall Evaporation Sunshine
##   <date>    <chr>      <dbl>  <dbl>   <dbl> <lgl>      <lgl>
## 1 2008-12-01 Albury      13.4   22.9     0.6 NA        NA
## 2 2008-12-02 Albury       7.4   25.1     0  NA        NA
## 3 2008-12-03 Albury      12.9   25.7     0  NA        NA
## 4 2008-12-04 Albury       9.2    28      0  NA        NA
## 5 2008-12-05 Albury      17.5   32.3     1  NA        NA
## 6 2008-12-06 Albury      14.6   29.7     0.2 NA        NA
## # ... with 16 more variables: WindGustDir <chr>, WindGustSpeed <dbl>,
## #   WindDir9am <chr>, WindDir3pm <chr>, WindSpeed9am <dbl>,
## #   WindSpeed3pm <dbl>, Humidity9am <dbl>, Humidity3pm <dbl>,
## #   Pressure9am <dbl>, Pressure3pm <dbl>, Cloud9am <dbl>, Cloud3pm <dbl>,
## #   Temp9am <dbl>, Temp3pm <dbl>, RainToday <chr>, RainTomorrow <chr>
```

```
dim(data)
```

```
## [1] 142193      23
```

I will check what percent of each column is NA

```
colSums(is.na(data))/nrow(data)*100
```

```
##      Date      Location      MinTemp      MaxTemp      Rainfall
## 0.0000000 0.0000000 0.4479827 0.2264528 0.9887969
## Evaporation      Sunshine      WindGustDir      WindGustSpeed      WindDir9am
## 98.7320051 98.1475881 6.5615044 6.5193083 7.0418375
## WindDir3pm      WindSpeed9am      WindSpeed3pm      Humidity9am      Humidity3pm
## 2.6569522 0.9480073 1.8495988 1.2476001 2.5388029
## Pressure9am      Pressure3pm      Cloud9am      Cloud3pm      Temp9am
## 9.8556188 9.8324109 37.7353316 40.1524688 0.6357556
## Temp3pm      RainToday      RainTomorrow
## 1.9171127 0.9887969 0.0000000
```

Evaporation, Sunshine, and both Cloud columns have too many NA's to be useful predictors so I can remove them

```
drops<-c("Evaporation","Sunshine","Cloud9am","Cloud3pm")
data<-data[,!(names(data)%in%drops)]
rm(drops)
na.exclude(data)%>%dim()
```

```
## [1] 112925      19
```

```
data<-na.exclude(data)
```

I will change RainTomorrow and RainToday from character class to factor so it is easier to sort them and use operators with them

```
data$RainTomorrow<-as.factor(data$RainTomorrow)
data$RainToday<-as.factor(data$RainToday)
nrow(data)
```

```
## [1] 112925
```

Over 100.000 observations is still plenty of data for predicting so I can remove all other NA's without impacting the results significantly Now I can split the data in test and training sets

```
set.seed(1, sample.kind = "Rounding")
test_index <- createDataPartition(y = data$RainTomorrow, times = 1, p = 0.1, list = FALSE)
train<- data[-test_index,]
temp <- data[test_index,]

validation <- temp %>%
  semi_join(train, by = "RainTomorrow") %>%
  semi_join(train, by = "RainToday") %>%
  semi_join(train, by = "Humidity3pm") %>%
  semi_join(train, by = "Pressure3pm") %>%
  semi_join(train, by = "Date")
```

Add rows removed from validation set back into data set

```
removed <- anti_join(temp, validation)
```

```
## Joining, by = c("Date", "Location", "MinTemp", "MaxTemp", "Rainfall", "WindGustDir", "WindGustSpeed")
```

```
data <- rbind(data, removed)
rm(url, removed, temp, test_index)
```

Data Exploration

Look at number of cities and date range of the data

```
range(train$Date)
```

```
## [1] "2007-11-01" "2017-06-25"
```

```
length(unique(train$Location))
```

```
## [1] 44
```


There is data from 44 Australian cities from 2007-2017

```
sort(table(train$Location),decreasing =TRUE)
```

```
##
##      Darwin      Hobart      Brisbane      Perth
##      2810      2754      2745      2714
## SydneyAirport MelbourneAirport PerthAirport MountGambier
##      2647      2639      2618      2615
##      Cairns      Mildura      Woomera      Townsville
##      2614      2600      2600      2596
## NorfolkIsland Ballarat      GoldCoast      Portland
##      2584      2568      2558      2542
##      Nuriootpa      Cobar      Wollongong      NorahHead
##      2517      2516      2512      2503
##      WaggaWagga      Adelaide      Canberra      Watsonia
##      2500      2479      2471      2462
##      Sale      AliceSprings      Bendigo      Moree
##      2456      2446      2428      2377
## CoffsHarbour      Walpole      PearceRAAF      Albury
##      2273      2256      2222      2179
## BadgerysCreek      Witchcliffe      Dartmoor      Sydney
##      2103      2078      2072      2060
##      Tuggeranong      Melbourne      Williamtown      Richmond
##      2046      2027      1960      1838
##      Launceston      Nhil      Uluru      Katherine
##      1394      1353      1289      611
```

Next I will figure out which data are the best predictors for Rain Tomorrow and remove the irrelevant columns

```
train%>%ggplot(aes(RainToday, fill= RainTomorrow))+
  geom_bar(position = "dodge")+
  scale_y_continuous(labels = function(x) paste0(round(x/nrow(train)*100,1),"%"),breaks = seq(0,nrow(tr
```




When it rains today, it is far more likely to rain tomorrow. In almost half of all occurrences when it Rained Today, it Rained Tomorrow.

```
monthly_rain_tomorrow<-train%>%filter(month(Date),RainTomorrow=="Yes")%>%count(month(Date))
monthly_days<-count(train,month(train$Date))
pct_rain_tomorrow_monthly<-data_frame("Month"=factor(month.abb,levels = month.abb),"Percent_Rain_Tomorrow"
```

Look at a histogram of the monthly average rainfall relative to the average rainfall which is represented by the horizontal line

```
pct_rain_tomorrow_monthly%>%ggplot(aes(Month, Percent_Rain_Tomorrow))+
  geom_col()+
  geom_hline(yintercept = mean(pct_rain_tomorrow_monthly$Percent_Rain_Tomorrow))+
  geom_text(aes(2,mean(pct_rain_tomorrow_monthly$Percent_Rain_Tomorrow), label= "Average", vjust=-.5))
```



The Australian rainy season is from June to August

Does the amount of rain today predict the likelihood of rain tomorrow? A few extreme days with lots of rain will make it difficult to visualize a density plot so by checking the standard deviation of days with rain (y) I can tell how to limit the x-axis


```
tail(sort(train$Rainfall),n=10)
```

```
## [1] 206.8 208.5 210.6 219.6 225.0 236.8 247.2 268.6 278.4 367.6
```

```
y<-filter(train,train$Rainfall>=0.1)
n<-filter(train,train$Rainfall==0)
sd(y$Rainfall)
```

```
## [1] 13.38233
```


```
ggplot(train,aes(Rainfall))+
  geom_density()+
  scale_x_continuous(limits = c(0.1,15))
```



The density plot shows that the majority of days with rain get under 13mm of rain


```
train %>%
  filter(Rainfall > 0.1) %>%
  ggplot(aes(x = Rainfall, fill = RainTomorrow, color = RainTomorrow)) +
  geom_histogram(aes(y = ((..count..) / sum(..count..))), position = "identity", alpha = 0.5) +
  scale_y_continuous(breaks = seq(0, 1, by = 0.02),
                     labels = scales::percent) +
  labs(x="Rainfall in mm", y="Percent of Rainy Days")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

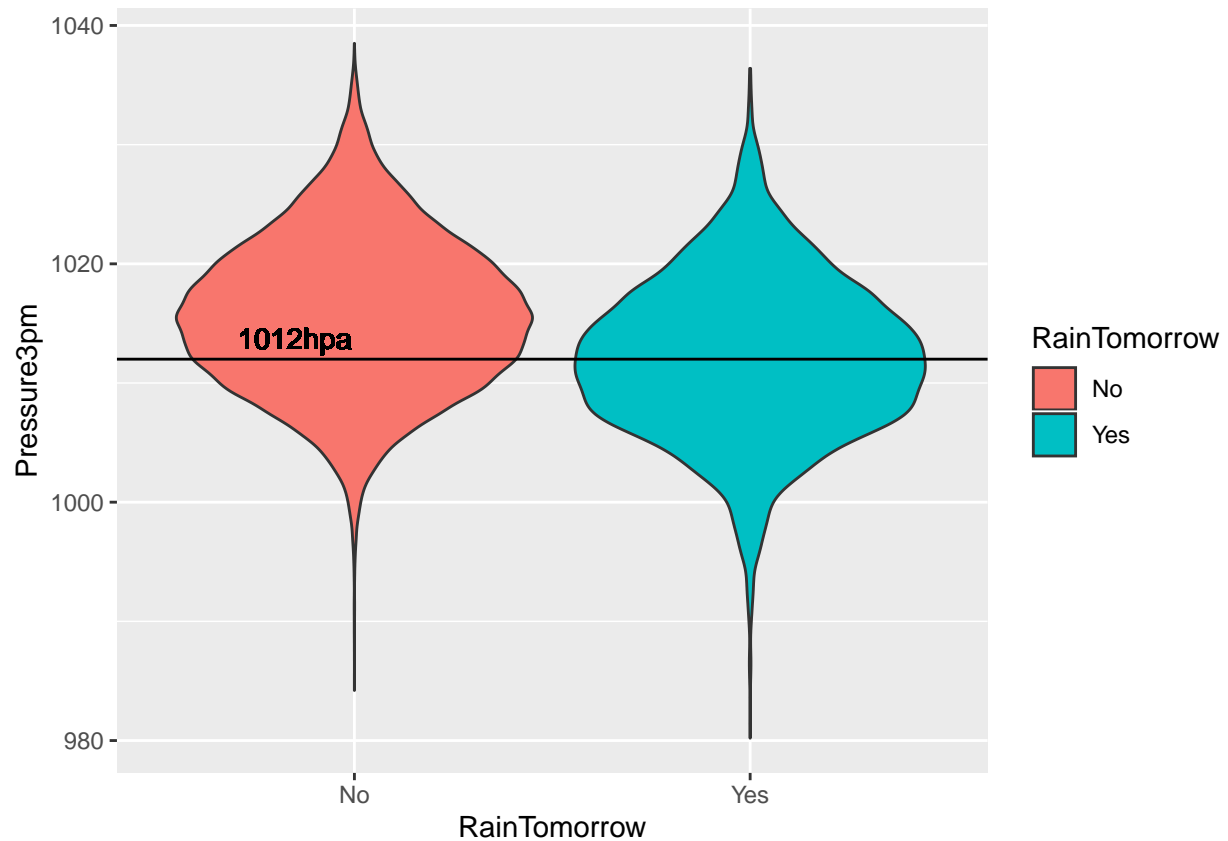
After a certain amount of rainfall, it is more likely to rain tomorrow, so I will zoom in to figure out the cutoff

```
train %>%
  filter(Rainfall > 1 & Rainfall < 15) %>%
  ggplot(aes(x = Rainfall, fill = RainTomorrow, color = RainTomorrow)) +
  geom_histogram(aes(y = ((..count..) / sum(..count..))), position = "identity", alpha = 0.5, binwidth = 1) +
  scale_x_continuous(breaks = seq(0, 15, by = 1)) +
  scale_y_continuous(breaks = seq(0, 1, by = 0.02),
    labels = scales::percent) +
  labs(x="Rainfall in mm", y="Percent of Rainy Days")
```



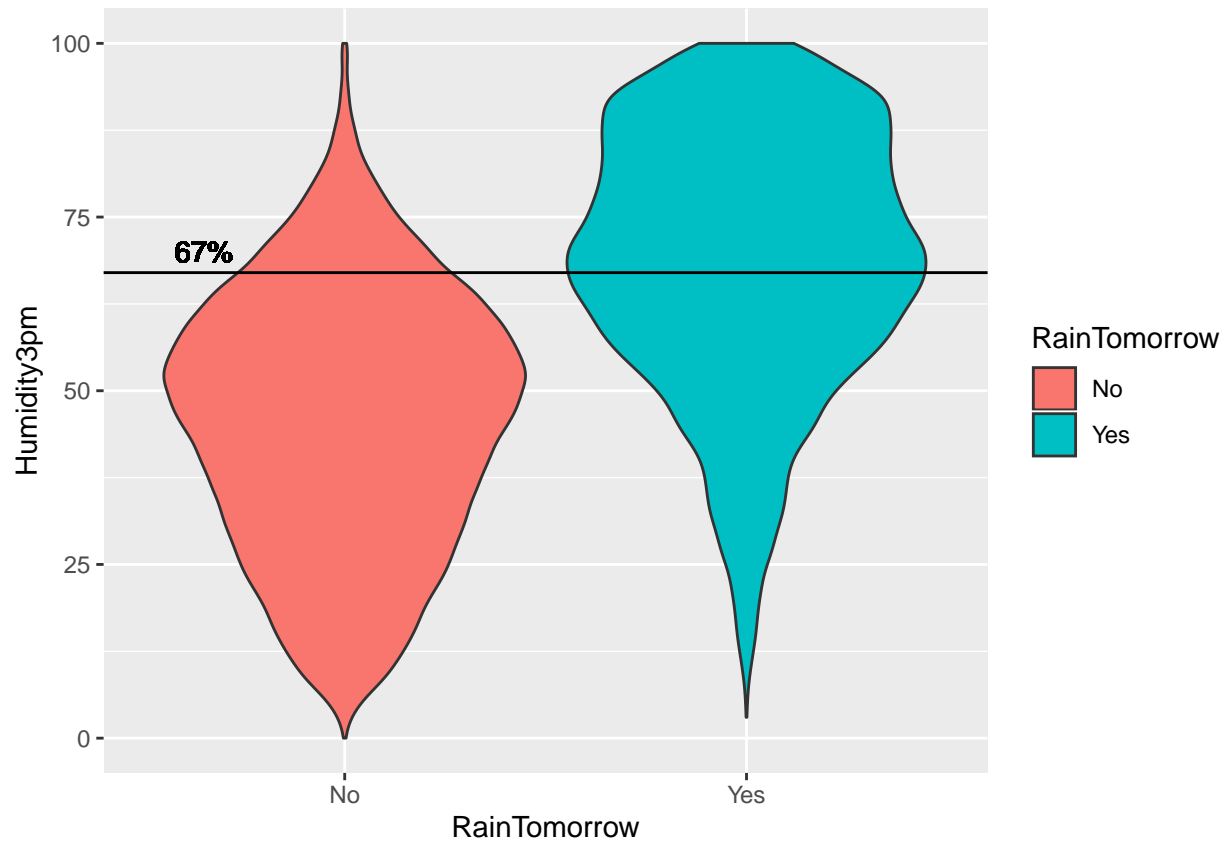
After 8mm of rain it is more likely to rain the next day than not rain Next I will check if Pressure at 3pm is a good predictor

```
n%>%ggplot(aes(x=RainTomorrow, y=Pressure3pm, fill= RainTomorrow)) +
  geom_violin()+
  geom_hline(yintercept = 1012)+
  geom_text(aes(.85,1012), label= "1012hpa", vjust=-.5)
```



As the violin chart shows, below 1012hpa, it is more likely to rain tomorrow Time to check if Humidity at 3pm is a good predictor of rainfall

```
train%>%ggplot(aes(x=RainTomorrow, y=Humidity3pm, fill= RainTomorrow)) +
  geom_violin()+
  geom_hline(yintercept = 67)+
  geom_text(aes(.65,67), label= "67%", vjust=-.5)
```



This shows us that it is more likely to rain tomorrow when humidity at 3pm is above 67%

Rain today, months June to August, more than 8mm of rainfall, pressure at 3pm below 1012hpa, and humidity at 3pm above 67% are the best predictors

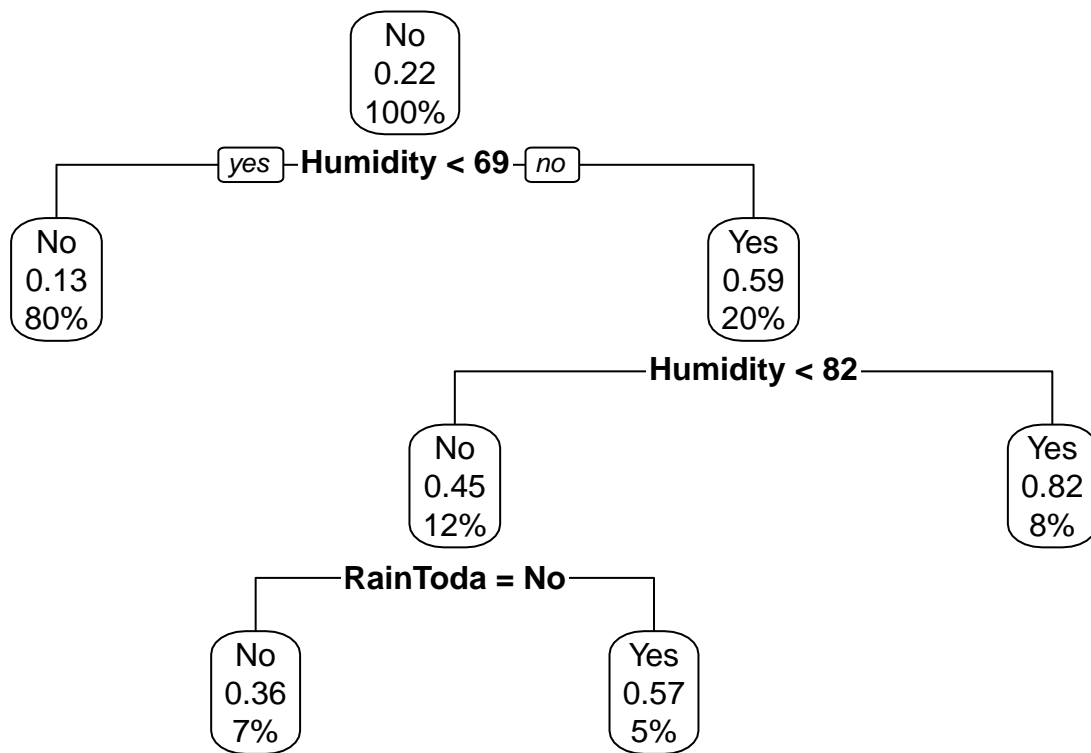
Methods/Analysis

Now to create the prediction model with these predictors

```
predictors<- c("Date","Rainfall", "Humidity3pm","Pressure3pm", "RainToday", "RainTomorrow")
train<-train[, (names(train)%in%predictors)]
Decision_Tree<-rpart(RainTomorrow~ RainToday + Date + Humidity3pm + Pressure3pm, data = train, method =
```

I can visualize the tree to see where it makes the breaks

```
prp(Decision_Tree, type = 2, extra = "auto", branch = 1)
```



The rpart function automatically chose the predictors Humidity and Rain Today

```
prediction_class_1<-predict(Decision_Tree, validation, type = "class")
confusionMatrix(prediction_class_1,validation$RainTomorrow,positive = "Yes")
```


```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   No  Yes
##           No 8369 1429
##           Yes  403 1065
##
##           Accuracy : 0.8374
##           95% CI : (0.8304, 0.8442)
##           No Information Rate : 0.7786
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.4469
##
##           McNemar's Test P-Value : < 2.2e-16
##
##           Sensitivity : 0.42702
##           Specificity : 0.95406
##           Pos Pred Value : 0.72548
##           Neg Pred Value : 0.85415
##           Prevalence : 0.22137
```

```
##          Detection Rate : 0.09453
##    Detection Prevalence : 0.13030
##      Balanced Accuracy : 0.69054
##
##      'Positive' Class : Yes
##
```

Using a confusionMatrix, I can check the metrics for success. The decision tree model has an 83.83% accuracy which is pretty good. However, the Sensitivity is 42.7% which is a bit low and the Balanced Accuracy is only 69.05%. But by using Random Forests I can improve the balanced accuracy and sensitivity

First I create a randomForest with the selected predictors and look at how many trees I need

```
rforest<-randomForest(RainTomorrow~ RainToday + Date + Humidity3pm + Pressure3pm, data = train, ntree=200)
plot(rforest)
```



The error stops changing significantly around 50 trees so I can limit the forest to 50 trees to save computing power

```
rforest<-randomForest(RainTomorrow~ RainToday + Date + Humidity3pm + Pressure3pm, data = train, ntree=50)
```

Now I can test to see if the randomForest improves the prediction of Rain Tomorrow

```
prediction_class_2<-predict(rforest,validation, type = "class")
confusionMatrix(prediction_class_2,validation$RainTomorrow, positive = "Yes")
```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 8215 1308
##           Yes  557 1186
##
##           Accuracy : 0.8345
##           95% CI : (0.8275, 0.8413)
##           No Information Rate : 0.7786
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.4618
##
## Mcnemar's Test P-Value : < 2.2e-16
##
##           Sensitivity : 0.4755
##           Specificity : 0.9365
##           Pos Pred Value : 0.6804
##           Neg Pred Value : 0.8626
##           Prevalence : 0.2214
##           Detection Rate : 0.1053
##           Detection Prevalence : 0.1547
##           Balanced Accuracy : 0.7060
##
##           'Positive' Class : Yes
##

```

With RandomForest, our new accuracy decreased slightly by 0.3% to 83.44%. However, the Sensitivity rose 4.7% to 47.47% and the Balanced Accuracy rose 1.52% to 70.57%

Conclusion

The Random Forest method slightly improved the balanced accuracy of the decision tree by raising the sensitivity. However, the Specificity dropped 1.74% which explains why the balanced accuracy did not increase further. With a specificity still above 90%, I think improving the low Sensitivity is worth the trade-off. One area of future research is altering the node size of the random forest to get a better accuracy, but my computer just kept freezing everytime I tried so with more computing power, perhaps a high accuracy could be achieved. Decision trees are good for making binary decisions, and with something as voluminous as weather data, there might be other machine learning models that could be better. One of the advantages of decision trees is it makes it easier to understand how the machine comes to a particular outcome since we can visualize each split of the branches.